

IT 18305 - DATABASE SYSTEMSCAT-3 (II<sup>nd</sup> Semester)

11.02.22

Part - A

③ Advantages for storing multiple relations in a single file:

- \* Complex structures can be implemented through the DBMS, thus increasing performance.

Disadvantages of storing multiple relations in a single file:

- \* Increases the size and complexity of the DBMS.

D For the two disk mirrored use, we assume A disk and B disk. In order to lose data, A and B need to be ~~failed~~ failed at the same time. If A is already failed and within 100000 hours B disk will fail, then data will be lost. The other case is B is already failed and within 100,000 hours A will fail, and then data will be lost.

If the mean time to failure of a single disk is 100000 hours and the mean time to repair is 10 hrs then the mean time to data loss is  $(100000 \times 2)^{1/2} \times 100$ .

## 5) Map database.

Map database management systems are software programs designed to efficiently store and recall spatial information.



- 2)  $B^+$  trees can be used to locate block containing the search key.

We choose  $B^+$  tree because it is a self-balancing tree data structure that maintains data and allows searches, sequential access, insertions and deletions.

## Part-B

### Q(b) Challenges in Maintaining Data Consistency

Data discrepancy occurs when the data in the target database deviates from the source database. The extent to which the data deviates depends on various factors.

Even when using products that replicate data reliably, such as Oracle, there remains potential causes of data discrepancy. ~~If the goal of the~~

#### Migration Errors

Different kinds of migration tools are employed to facilitate the initial load of the target databases before replication can begin. Differences in configuration for handling data by the migration tools and replication products can ~~not~~ result in data discrepancies.

#### Lift & Shift workload to Cloud

Since the world is moving towards the cloud, the lift & shift of database from on-premises to cloud is the need of today's IT world.

#### Differences in Source and Target

Diff in source and target database configuration, for example different encodings, locales, endianness, or database versions can cause subtle discrepancies to happen during migration and replications.

## Instantiation Errors

Before migration or replication can begin, the target databases will need to be instantiated with the correct schema and constraints.

## Configuration Errors

Improper and unintended configuration of replication products can cause discrepancies.

## Gaps in Replication

Although replication is enabled between source and target databases and is working perfectly well, there are instances where data inserted on the source will not be replicated.

## Replication Latency

With asynchronous replication, there will be a short lag between changes to the source database and delivery of those changes to the target.

## User errors

Often target databases are created to offload query processing from the source database. This enables rich operational reporting without impacting the application's running.

## Application errors

Applications that use target databases can potentially change data due to faulty logic as well as application upgrades.

The technology requirements for managing data consistency are:

- \* High speed, low impact data comparisons.
- \* Support for heterogeneous databases.
- \* Capability for handling large data volumes.
- \* Flexible options for managing data comparisons.
- \* Support for live databases with constantly changing data.
- \* Minimally intrusive.
- \* Comparison of only changed data in continuous replication.
- \* Comparison of huge table through automated & manual partitions.
- \* Data comparison reports for auditing purposes.
- \* Zero downtime of source and target systems.
- \* Capability to identify data ~~inconsistency~~ inconsistency.
- \* Low impact on hardware and network resources.
- \* Flexible reporting for varying L roles and access levels.
- \* Data security.
- \* Easy to use, understand, configure, deploy and integrate.

### Part - C

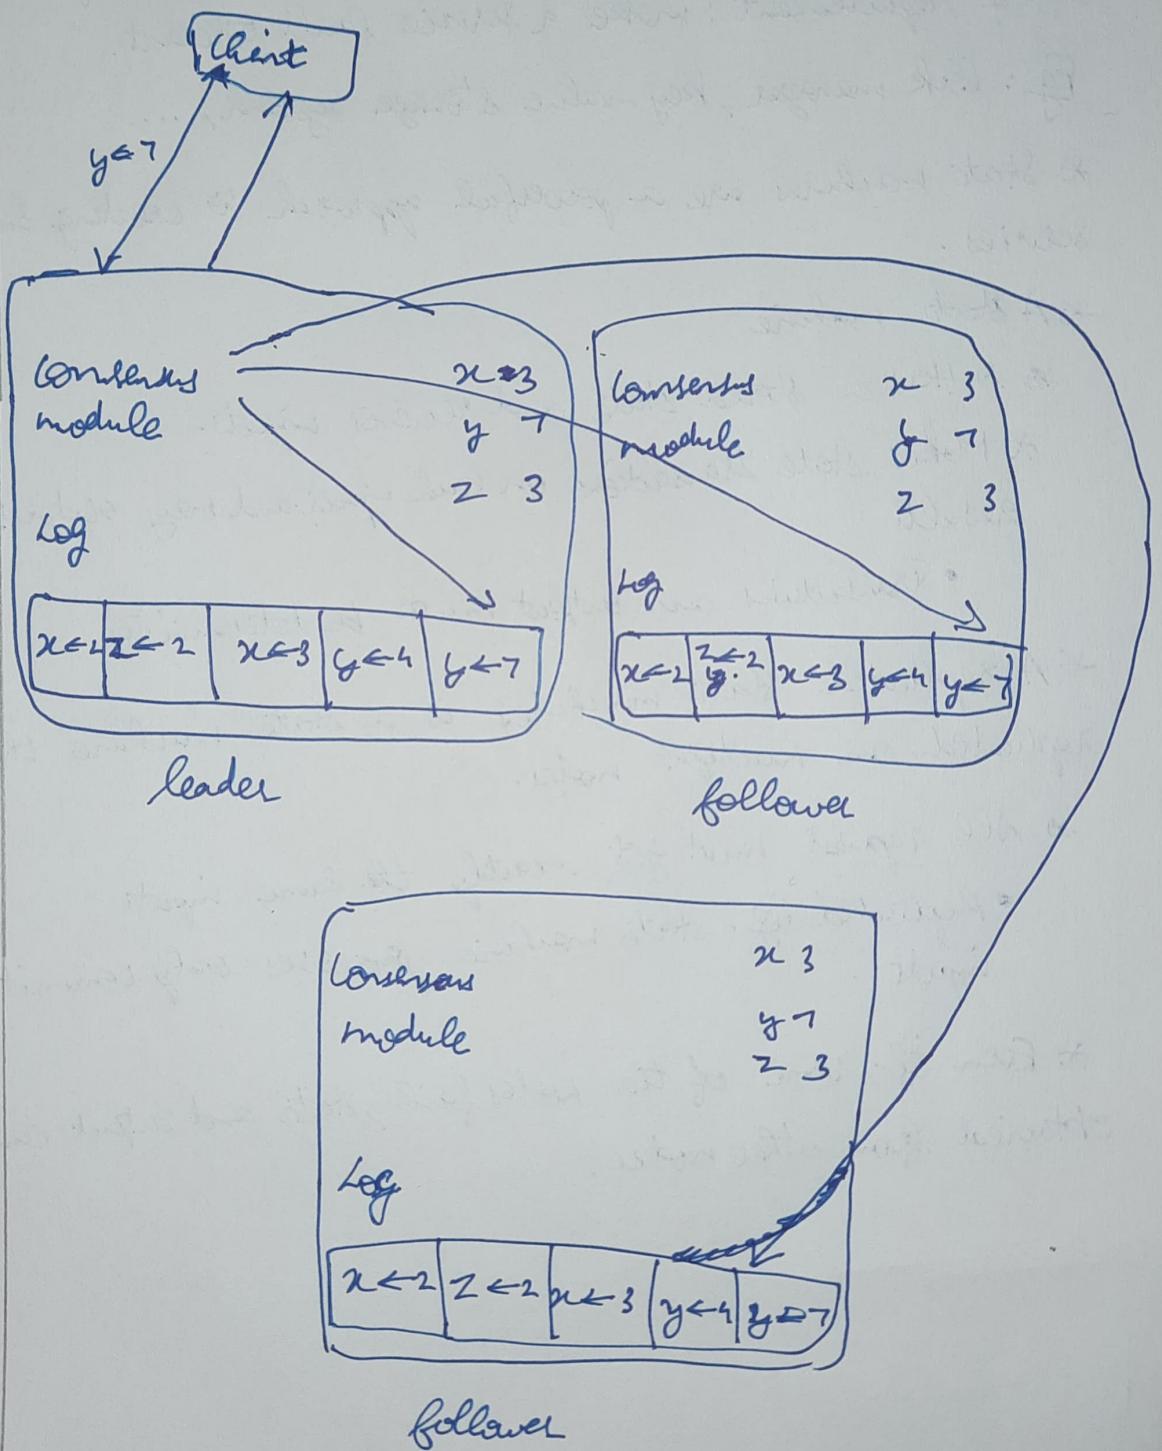
D(a)

#### Fault Tolerant Services using Replicated State Machines

- \* Key requirement: make a service fault tolerant.  
Eg: lock manager, key-value storage system, ....
- \* State machines are a powerful approach to creating such services.
- \* A state machine
  - \* Has a stored state, and receives inputs.
  - \* Makes state transitions on each input, and may output some results
    - Transaction and output must be deterministic
- \* A replicated state machine is a state machine that is replicated on multiple nodes.
  - \* All replicas must get exactly the same inputs
    - Replicated log. State machine processes only committed inputs.
  - \* Even if some of the nodes fail, state and output can be obtained from other nodes.

## Replicated State Machine

- \* Replicated state machine based on replicated log.
- \* Example commands assign values to variables.



Leader dables log record committed after it is replicated as a majority of nodes.

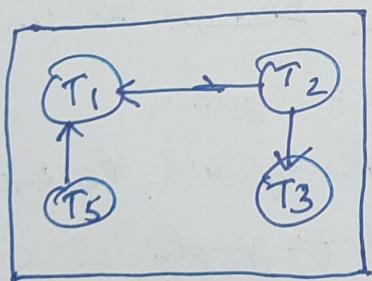
## Uses of Replicated State Machines

- \* Replicated state machines can be used to implement wide variety of services.
  - \* Inputs can specify operations with parameters.
  - \* But operations ~~to~~ must be deterministic.
  - \* Result of operation can be sent to from any replica.
    - Gets ~~not~~ executed only when log record is committed in replicated log.
- \* Example : fault-tolerant lock manager.
  - \* State : lock table.
  - \* Operations : lock requests and lock releases.
  - \* Output : grant, or rollback requests on deadlock.
  - \* Centralized implementation is made fault tolerant by simply running it on a replicated state machine.
- \* Fault tolerant key value-store
  - \* State : key value storage state
  - \* Operations : get() and put() are first logged.
- \* Google spanner ~~uses~~ uses replicated state machine to implement key value store.
  - \* Data is partitioned, and each partition is replicated across multiple nodes.

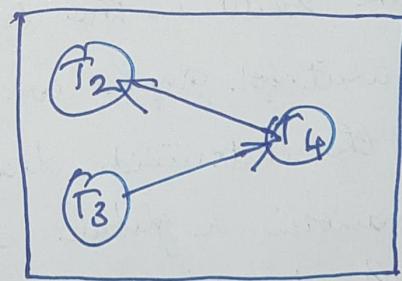
## 8(b) Deadlock Handling

The deadlock handling prevention and deadlock-detection algorithms can be used in a distributed system, provided that modifications are made. For example, we can use the tree protocol by defining a global tree among the system data items. Similarly, the timestamp-ordering approach could be directly applied to a distributed environment.

Deadlock prevention may result in unnecessary waiting and rollback. Furthermore, certain deadlock-prevention techniques may require more sites to be involved in the execution of a transaction than would otherwise be the case. If we allow deadlocks to occur and rely on deadlock detection, the main problem in a distributed system is deciding how to maintain the wait-for graph.



site S<sub>1</sub>



site S<sub>2</sub>

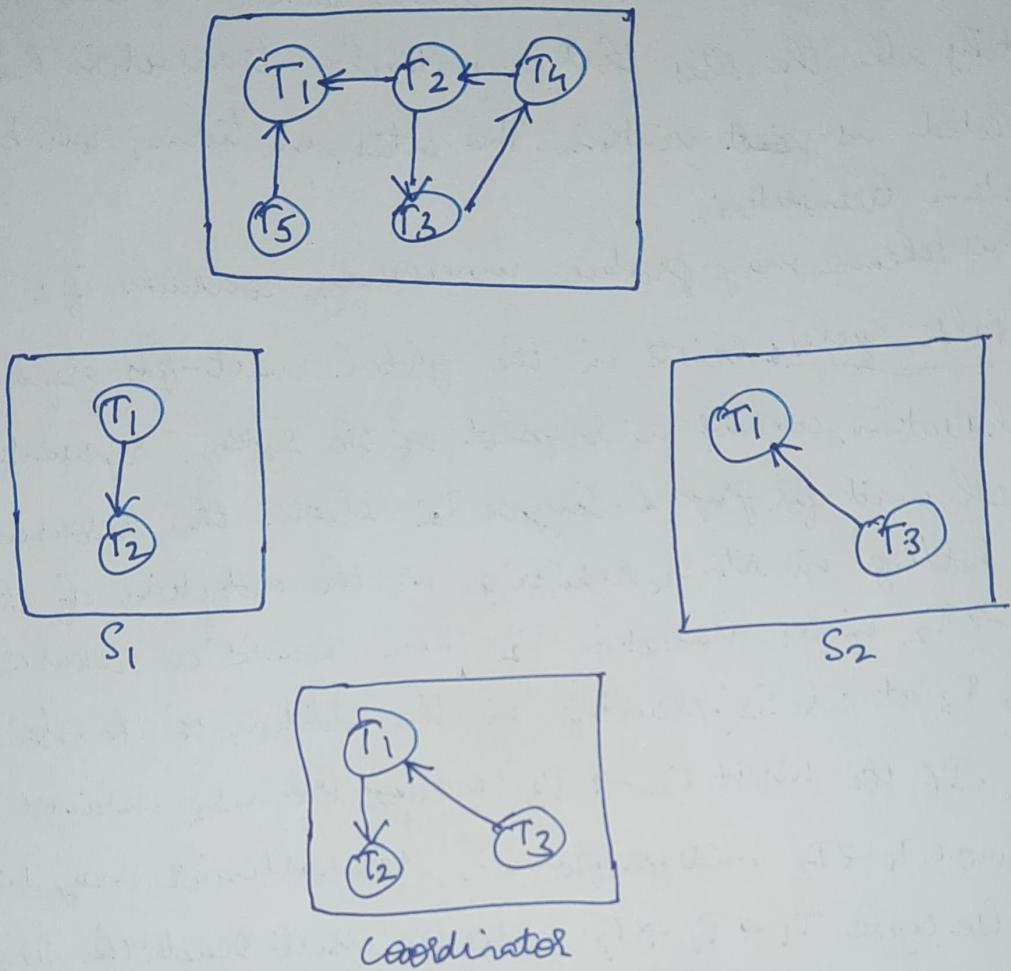
Common techniques for dealing with the issue require that each site keep a local wait-for graph. The nodes of the graph correspond to all the transactions (local as well as non-local) that are currently either holding or requesting any of the items local to that site. For example,

The above diagram depicts a system consisting of two sites, each maintaining its local wait-for graph. Note that transactions  $T_2$  and  $T_3$  appear in both graphs, indicating that the transactions have requested items at both sites.

These local wait-for graphs are constructed in the usual manner for local transactions and data items. When a transaction  $T_i$  on site  $S_1$  needs a resource in site  $S_2$ , it sends a request message to site  $S_2$ . If the resource is held by transaction  $T_j$ , the system inserts an edge  $T_i \rightarrow T_j$  in the local wait-for graph of site  $S_2$ . Clearly, if any local wait-for graph has a cycle, deadlock has occurred. On the other hand, the fact that there are no cycles in any of the local wait-for graphs does not mean that there are no cycles in any of the local wait-for graphs. It does not mean there are no deadlocks. Each wait-for graph is auxiliary; nevertheless, a deadlock exists in the system because the union of the local wait-for graphs contains a cycle.

In the centralized deadlock approach, the system constructs and maintains a global wait-for graph in a single site: the deadlock-detection coordinate. Since there is communication delay in the system, we must distinguish between the two types of wait-for graphs. The real graph describes the real but unknown state of the system at any instance in time, as would be seen by an omniscient observer. The constructed graph is an

approximation generated by the controller during the execution of the controller's algorithm.



False cycles in the global wait-for graph

The global wait-for graph can be reconstructed & updated under these conditions:-

- \* Whenever a new edge is listed in or removed from one of the local wait-for graphs.
- \* Periodically, when a number of changes have occurred in a local wait-for graph.
- \* whenever the coordinator needs to invoke the cycle-detection algorithm.

When the coordinator invokes the deadlock - detection algorithm, it searches its global graph. If it finds a cycle, it selects a victim to be rolled back. The coordinator must notify all the sites that a particular transaction has been selected as ~~victim~~ victim. The sites, in turn, roll back the victim transaction.

This scheme may produce unnecessary rollbacks if:

\* False cycles exists in the global wait-for graph. As an illustration, consider a snapshot of the system represented by the local wait-for graphs. Suppose  $T_2$  releases the resource that it is holding in site  $S_1$ , resulting in the detection of the edge  $T_1 \rightarrow T_2$  in  $S_1$ . Transaction  $T_2$  then requests a resource held by  $T_3$  at site  $S_2$ , resulting in the addition of the edge  $T_2 \rightarrow T_3$  in  $S_2$ . If the insert  $T_2 \rightarrow T_3$  message from  $S_2$  arrives before the remove  $T_1 \rightarrow T_2$  message from  $S_1$ , the coordinator may discover the false cycle  $T_1 \rightarrow T_2 \rightarrow T_3$  after the insert. Deadlock recovery may be initiated.

\* A deadlock has indeed occurred and a victim has been picked, while one of the transactions was aborted for reasons unrelated to the deadlock.